


A RESEARCH AND LEARNING LAB INCUBATED AT THE AGENCY FUND

The **Agency** & AI Lab

A²L is a decentralized, virtual, and open lab with two jobs: **design** AI that strengthens human agency in the social sector, and **measure** what AI does to the agency of the people who use it.

An interactive map of organizations and their AI products in TAF's accelerator

INTERACTIVE Click the map, chips, or legend to focus an org · click empty space to clear, or Reset



DG **Digital Green**
2025 COHORT

FarmerChat

A multilingual AI agronomy advisor for smallholder farmers, using text, voice, and images to deliver localized recommendations.

L1 · MODEL Model stack is live across 16 languages, with text, voice, image evals and over 100% improvement over the base model.

L2 · PRODUCT 1M+ users have generated 8M+ queries; 30% are monthly active and average 8.7 queries per month.

L3 · USER About 70% of active users report applying FarmerChat advice; 80% share advice with peers.

L4 · IMPACT Kenya assessment finds adoption of cost-saving practices; income and climate-resilience outcomes are in ongoing RCTs.

AI PRODUCT: INDIA · KENYA · NIGERIA · ETHIOPIA · BRAZIL
ORGANIZATION OPERATES: INDIA · KENYA · NIGERIA · ETHIOPIA · BRAZIL

Working with these organizations surfaced **two gaps related to human agency** in AI product work

OBSERVATION FROM THE ACCELERATOR

AI tools are advancing service delivery faster than the field is learning what they do to people's agency.



THE DESIGN GAP

AI products are built to optimize service delivery, not explicitly to strengthen people's agency.

Despite AI's potential to strengthen human agency, NGO products often treat agency as a byproduct rather than a design objective.

Decades of behavioral science research identify agency as a key mechanism for improving development outcomes, yet it is rarely **built into AI design, testing, or iteration.**



THE MEASUREMENT GAP

AI products are underused as windows into how agency evolves.

Many products may not generate as many signals of agency as tools like ChatGPT or Claude Code do. When signals exist, they often come from power users, and the field lacks a shared standard or analytical tools for **finding, gathering, and interpreting agency** inside human-AI traces.

Same mother, same question — two very different conversations

INFORMATION DELIVERY

<

A

Amma Health Line

AI assistant · online

Today

hello. my feet r swelling in the evening. i am 7 months pregnant

10:14 ✓

Hello Priya. Mild evening swelling is common in month 7. It helps to:
1. Rest with your feet raised.
2. Drink water through the day.
3. Avoid standing for long.
Go to the clinic if swelling is sudden or comes with headache or blurred vision.

10:15

ok. can i still do house work

10:22 ✓

Light housework is fine. Avoid heavy lifting and take short breaks.

10:23

ok thank you

10:26 ✓

Message

AGENCY DEVELOPMENT

<

A

Amma Health Line

AI assistant · online

Today

hello. my feet r swelling in the evening. i am 7 months pregnant

10:14 ✓

Hello Priya. Evening swelling is common in month 7, and resting with feet raised helps. Have you managed your ANC check-up this month?

10:15

not yet. clinic is far and my husband takes the scooter for work

10:21 ✓

That is a real constraint. The 7th-month check matters — they will look at this swelling and your BP. Which day is he home?

10:22

thursday, home by 2

10:24 ✓

Then a plan: Thursday at 3 he drops you at the clinic. Carry your mother-child card. Shall I remind you at 1?

10:25

yes pls. thursday 3 pm i will go

10:27 ✓

Message

LEFT PHONE

The answer is **correct** — and the conversation ends where it began. Nothing changed in what Priya will do next.


RIGHT PHONE

The same opening becomes **a goal she owns**:

- Open question surfaces the real barrier
- Concrete if-then plan around her constraints
- Commitment stated in her own words, plus a reminder

Same model, same facts. The difference is **a design choice** — and social-psychological research tells us which conversational moves build agency.

Our worldview of agency in a flywheel



The loop can turn in either a good or a bad way.

When it turns well, capable choices bring results that promote a positive cycle. When it turns badly, failed choices result in a negative cycle.

AI products sit inside it, at every point

An AI tool can push any of the four elements in either direction, while generating lots of digital behavioral traces.

That's exactly why we need to design AI for human agency and be able to concretely measure agency.

Studies show AI can **expand** agency, and also hamper it

THE FULL FINDING

CAN EXPAND

In a higher-resource AI use setting with advanced users, people stayed with tools that **made them feel more capable**, even when the tools made mistakes. The study points to five ways AI can expand agency: instrumental, cognitive, relational, and structural.

FLIP BACK

THE FULL FINDING

CAN HAMPER

A ChatGPT-style math tutor boosted practice scores by 48%, but students scored **17% below the control group** once the AI was removed. A quarantined tutor that offered hints rather than solutions eliminated the harm, yet students using either tool remained unaware of the gap between feeling supported and actually learning.

FLIP BACK

THE FULL FINDING

CAN HAMPER

Across 1,325 participants, an accurate chatbot raised reasoning accuracy by 25 percentage points, but a faulty one drove it 15 points below unaided performance. People followed wrong AI advice on **roughly four in five consulted trials** and became more confident even when misled; time pressure, incentives, and feedback reduced but did not eliminate the effect.

FLIP BACK

These are evidence-focused studies from higher-resource settings where there are enough engagement signals to study agency. They offer early clues for social sector AI: **how design can expand or hamper agency**, and **how agency can be measured directly rather than inferred from usage alone**.

Sources: [How AI Is Affecting Human Agency](#) (Jigsaw, July 2025) · [Generative AI without guardrails can harm learning](#) (PNAS 122(26), 2025) · [Thinking, Fast, Slow, and Artificial: the rise of cognitive surrender](#) (Wharton working paper, 2026).

PART 2 · LEARNING PROJECTS

Repeatable motions the field can reuse.

What the lab has shipped against the two gaps so far: assistants, papers, and skills that put the agency literature to work — open and reusable by any organization.

Three bodies of work, aimed squarely at the two gaps

What the lab has built so far

1	LEARNING PROJECTS • REPEATABLE MOTIONS Put the agency literature in practitioners' hands Assistants, papers, and skills any organization can pick up and reuse. NotebookLM research assistants Evaluating AI Across Cultures (PNAS Nexus) Behavioral-science skills with Irrational Labs	ADDRESSES The design gap The measurement gap
2	RESEARCH PROJECTS • DESIGNING AI FOR AGENCY Study what agency-centered design does inside real products Field research with NGO partners, with the psychology brought in from the start. Noora Health copilot study Wise interventions × AI design agenda	ADDRESSES The design gap The measurement gap
3	RESEARCH PROJECTS • MEASURING AGENCY Build instruments that read agency from real usage A field-tested method, and a flagship study to validate it at scale. Nairobi rubric sprint Penn LLM-judge study with six NGOs	ADDRESSES The design gap The measurement gap

Also shipped: an Agency/AI University course demo and an agency-based intervention playbook (in progress) — everything open and reusable.

A research library that answers back

Learning projects · Google NotebookLM assistants built at The Agency Fund

● ● ● Agency Intervention Learning Assistant

SOURCES

📄 Wise interventions: a review (Walton & Wilson)

📄 Values affirmation in health settings

📄 Goal setting & implementation intentions

📄 Social norms & belonging interventions

📄 + TAF's curated agency research library

We run a maternal-health program in Bihar. Which agency interventions have worked in health contexts, and what would adapting one take?

Three intervention families have strong evidence in health settings: **implementation-intention prompts** for care-seeking 1, **values affirmation** before threatening health information 2, and **norm-based framing** of new practices 3. For a WhatsApp program, the closest analog is the if-then planning protocol in 1 — adaptation would mean localizing the barriers list and pre-testing the framing with your field team.

Ask about your sources...

➤

Agency Intervention Learning Assistant

Answers intervention-design questions from TAF's curated agency research library — grounded, cited, and open to any NGO team.

Storytelling design workshop

The same motion, applied: a NotebookLM grounded in narrative psychology helped **Noora Health** teams design a storytelling intervention.

The literature starts working **for the team** — decades of research, one question away, instead of behind paywalls and jargon.

Every benchmark and metric carries **cultural defaults** — we mapped where they hide

Learning projects · multi-author paper from the accelerator's working group, PNAS Nexus (2026)

PNAS Nexus

2026

Evaluating AI Across Cultures: A Multilevel Framework for Generative AI in Global Development

Culture is not a variable to control for. It operates as a set of **defaults** — invisible assumptions about what "good" looks like — baked into every evaluation tool.

A WORKING GROUP ACROSS

The Agency Fund · Stanford · Microsoft Research · Center for Global Development · Carnegie Mellon · Michigan State · UT Dallas · George Washington

LEVEL 1
Model

Benchmarks test **Western academic knowledge in English** — and say nothing about care advice in the Rift Valley.

LEVEL 2
Product

Engagement metrics inherit **commercial defaults**: attention is not utility, and more usage is not more welfare.

LEVEL 3
User

Psychological measures assume **WEIRD constructs and survey norms** that do not travel across cultures.

LEVEL 4
Impact

Trials chase a **moving target**: model updates and prompt changes silently shift the product mid-study.

When the tool's defaults don't match the deployment context, evaluation produces **artifacts, not evidence**.

Behavioral science that **rides along** inside everyday AI work

Learning projects · Claude Skills co-developed with Irrational Labs

● ● ● an NGO product team's AI assistant

> use the **behavioral-science** skill to audit our chatbot onboarding flow

Running the behavioral diagnosis...

- ✓ **Friction:** five setup questions before any value — front-load one quick win instead.
- ✓ **Motivation:** benefits framed for the org, not the user — reframe around her goal.
- ✓ **Timing:** the commitment ask lands before trust is built — move it after first success.

> draft the revised first screen with one quick win

Drafting...

- ✓ Two variants ready to A/B test — each grounded in the skill's checklist and cited to the underlying studies.

Expertise, packaged as skills

Irrational Labs' frameworks — diagnostics, checklists, intervention patterns — encoded as skills any AI assistant can load.

A repeatable motion

What used to require a consulting engagement now travels with the team's own tools — and improves with every co-developed release.

The same motion works for any body of expertise:
encode it once, reuse it everywhere the field works.

Open source: github.com/agency-fund/taf-skills · [behavioural-skills demo](#)

PART 3 · RESEARCH: DESIGNING AI FOR AGENCY

Design the agency in — don't bolt it on.

Field research inside real NGO products, asking what AI design does to users' agency — with the psychology brought in from the start.

A live copilot at scale: what does AI do to the humans in the loop?

Research · with Noora Health and the Stanford Center for Health Education

THE SETTING

Noora Health's **Mobile Care Companion** answers families' health questions over WhatsApp. Since November 2025, a **GenAI copilot** drafts language-matched replies for frontline support staff — who review, edit, and send every message.

240_k

queries per month across India

80_k

users reached monthly

15_k

health facilities

6

languages, Hindi to Odia

THE STUDY · MIXED METHODS

Copilot access

9 staff

randomized 1:1

No copilot

9 staff

RQ1

How does the copilot change staff **experience** — trust, agency, everyday practice? Interviews, plus every edit they make to an AI draft as a behavioral trace.

RQ2

Do **conversations** measurably differ? Expert rubrics scaled with a validated LLM judge across the full corpus, in English and Indic languages.

Then

Findings **triangulated**: multilevel models over conversation metrics, interpreted alongside what staff say in interviews.

The bigger arc: Noora is designing **personalization** into its AI products — this study builds the evidence base for doing it agency-first.

The psychology is ready — it's just **not in the room** where AI gets designed

Research · the design agenda ahead

Wise interventions

Small, precise shifts in how people make meaning — with outsized, lasting effects on outcomes.

Learning sciences

How people actually build skill and understanding — not just how they consume content.

Behavioral science

Diagnosing the frictions and motivations that decide whether advice becomes action.

Social psychology

Identity, belonging, and norms — present in every conversation, designed for in almost none.

Nature Reviews Psychology

2026

Psychologists must be involved in building conversational AI chatbots

Xuan Zhao · Stanford University

Chatbots could democratize personalized support — or deepen isolation and risk. The difference is whether **psychological theory and methods are embedded** in design, evaluation, and monitoring.

[nature.com/articles/s44159-026-00564-z](https://www.nature.com/articles/s44159-026-00564-z)

Most NGO AI products in the Global South are built without any of this expertise. That is the design gap — and **the lab's opening**.

PART 4 · MEASURING AGENCY

You cannot design for what you **cannot see.**

Agency measurement is interesting, necessary, and complex. The next four slides show the approach we teach and use.

What you need to measure also depends on the forms of agency and how agency manifests in your product.

Hypothetical example: an AI mentor helping learners develop a viable business idea

Agency with respect to the end goals

Can users make progress toward outcomes they value, whether through, around, or beyond the tool?

Agency with respect to the AI tool

Can users understand, question, adapt, and control the tool, deciding when and how to rely on it?

Internally experienced agency

People's thoughts and feelings

ON-PLATFORM DATA • IN-APP MICRO-SURVEY
Short in-app surveys providing users' general sense of self-efficacy or perceived control over their ultimate goal over time (e.g., "How confident is the user about developing a successful business idea? How much control do they feel over reaching their ultimate goal?")

OFF-PLATFORM DATA • PHONE OR IN-PERSON SURVEY
Longer phone/in-person surveys using validated agency measures that systematically probe users' general sense of self-efficacy or perceived control over their ultimate goal over time.

FLIP BACK

Externally expressed agency

People's choices and actions

ON-PLATFORM DATA • CONVERSATION LOGS
Longitudinal user signals of agency in AI conversation logs, especially those related to setting and navigating longer-term goals or demonstrating resilience despite difficulties.

OFF-PLATFORM DATA • RUBRIC-RATED PITCHES
Longer phone/in-person surveys tracking how users develop their business ideas; rubric-rated observations of ideas and pitches; or classroom observation of teamwork, ideally linked to downstream outcomes like business earnings.

FLIP BACK

ON-PLATFORM DATA • IN-APP MICRO-SURVEY
Short in-app surveys providing whether users felt able to understand, question, or adapt an AI mentor's advice.

OFF-PLATFORM DATA • VALIDATED SCALES
Longer phone/in-person surveys using validated agency measures that systematically probe users' trust in IA and perceived dependence versus empowerment when using a tool.

FLIP BACK

ON-PLATFORM DATA • DIGITAL TRACES
Digital traces (e.g., interaction patterns or conversation logs) showing whether users understood, question, or adapt an AI mentor's advice.


OFF-PLATFORM DATA • OBSERVATION AND INTERVIEWS
Observation or interview data (e.g., usability studies) on tool usage for developing business ideas.

FLIP BACK

 **Mind the social setting:** norms and relationships can also shape these questions; interpret the answers in your users' contexts.

Agency measurement sits at **Level 3**

The four-level evaluation framework · eval.playbook.org.ai



Nobody can run an RCT on every product change. **Validated agency measures at Level 3** move when the product changes, and they speak to what Level 4 cares about. That also allows teams to anchor their A/B tests on intermediate measures that are psychologically and behaviorally meaningful and relatively easy to collect.

Teams should specify what agency means for **their specific users and contexts**

1

Define agency for your local problem

What kind of agency is the tool supposed to support? **For whom? To what end?** Anchor the definition in your program's theory of change.



2

Measure with the right tools, on the right data, at the right moment


Combine **on-platform** signals, such as in-app surveys and digital traces, with **off-platform** surveys, interviews, observations, and records.



3

Test and validate your measurement tools iteratively

Test for evidence of **reliability and validity** for your context.



We put the method **in a room**

Field-tested · Nairobi onsite · July 2026

- 1** Teams of grantees, funders, and staff coded **30 synthetic AI-assistant conversations** from education, health, and agriculture.
- 2** Everyone labeled the same **binary rubric**, built on the agency flywheel: motivation, action and choice, links to real-world outcomes, and updated mental models.
- 3** **Calibrate** showed people-vs-people agreement live. Then an **LLM judge** ran the exact same rubric on the same conversations.

What the room learned

Binary questions force choices on borderline cases, and those forced choices surface the definitional debates that actually matter for measurement.

Once a rubric earns agreement between people, an AI can apply it consistently across **thousands of conversations**. That is what makes agency measurement scale.

PART 5 · THE FLAGSHIP STUDY

Now we are building the instrument.

A joint research program with the University of Pennsylvania, running across six NGO platforms in five countries.

An **LLM judge** for human agency

Hamsa Bastani

The Wharton School, University of Pennsylvania

Osbert Bastani

Computer and Information Science, University of Pennsylvania

Zezen (Michael) Wu

The Agency Fund

The same Penn team ran the PNAS 2025 study on AI guardrails and learning.

Digital platforms already produce conversations, decisions, and, where partners can link them, real outcomes. We design **LLM judges** that read those transcripts against a structured rubric and infer each component of agency, with scores that **admit uncertainty**.

Intentionality

Does the user form plans that serve their goals?

Forethought

Do they weigh what would happen under different courses of action?

Self-reactiveness


Do they carry plans out and adapt when new information arrives?

Self-reflectiveness


Do they look back on outcomes to improve the next decision?

Framework grounded in Social Cognitive Theory; it extends to other definitions of agency, including the capability approach.

From raw transcripts to trusted scores



WHAT MAKES THIS HARD, AND HOW WE HANDLE IT



Six partners, three domains, five countries

HEALTH



Jacaranda Health · Kenya

PROMPTS, maternal-health SMS with AI triage.

3.8M+

mothers reached with 22 county governments

EDUCATION



Nova Escola · Brazil

WhatsApp AI lesson planner on the national curriculum.

150K teachers in year one



The Apprentice Project · India

TAP Buddy project-based learning for ages 9 to 17, now adding mother-tongue voice AI.



Udhyam Learning Foundation

· India

AI help bot inside a real-business entrepreneurship program.

LIVELIHOODS



Digital Green

· Kenya, Ethiopia, India

FarmerChat agronomy advisor.

1.3M farmers



Kabakoo Academies · Mali

WhatsApp AI mentor with video reflections, often in Bambara.

Every partner already runs an AI product with real users and real stakes. Together they give the study conversations in **text, voice, and video** across languages and cultures.

What each partner brings, and what we will learn together 1 of 2



DIGITAL GREEN
KENYA · ETHIOPIA · INDIA

PRODUCT

FarmerChat answers smallholders' agronomy questions; version 2 adds visual learning cards and profile-based personalization.

DATA

1.3 million users and hundreds of thousands of conversations, linkable to adoption and yield surveys from ongoing RCTs.

STUDY

Find agency signals among power users, track how they shift from version 1 to version 2, and test whether they predict real adoption, alongside the World Bank MAGNET agency measures Digital Green already uses.



JACARANDA HEALTH
KENYA

PRODUCT

PROMPTS supports new and expecting mothers over two-way SMS with answers, AI risk triage, and referrals into care.

DATA

Conversation logs from a service that has reached **3.8M+ mothers**, linkable to routine surveys and care-visit records.

STUDY

Read how mothers navigate care pathways, link those signals to visits, and design proactive nudges for mothers who do not yet know what to ask.



KABAKOO ACADEMIES
MALI

PRODUCT

A **WhatsApp AI mentor** guides youth upskilling; learners answer curriculum prompts with short reflection videos, often in Bambara.

DATA

About **5,000 learners, 15,000 reflections**, and 4,000 video transcripts. A waitlist RCT already showed gains of 0.14 to 0.25 SD in growth mindset.

STUDY

Mine reflections for agency signals and connect them to income and community standing years later, under strict privacy protections for at-risk youth.

What each partner brings, and what we will learn together 2 of 2



NOVA ESCOLA
BRAZIL

PRODUCT

A **WhatsApp lesson planner** built on 40 years of teacher content and Brazil's national curriculum, made for low-connectivity classrooms.

DATA

150,000 teachers generated half a million lesson plans in the first year, and 76,127 in February 2026 alone.

STUDY

Does co-creating plans in conversation build lasting teaching practice, or only save time? Vary copilot features and measure what teachers learn.



THE APPRENTICE PROJECT
INDIA

PRODUCT

TAP Buddy guides students ages 9 to 17 through hands-on projects in coding, art, science, and financial literacy over WhatsApp; a new voice AI reaches parents and students in Hindi, Marathi, and Punjabi.

DATA

Multimodal project submissions (text, images, video) plus the new voice transcripts from family calls.

STUDY

Measure creativity and 21st-century skills as agency signals, and learn how bot-delivered feedback changes what students attempt next.



**UDHYAM LEARNING
FOUNDATION**
INDIA

PRODUCT

Students form teams, receive seed money, run **real businesses**, and pitch on video; an AI help bot answers program and learning questions around the clock.

DATA

Help-bot chats, idea submissions, pitch videos, and sales records, with outcome data on about **2,000 students**, plus a McGill-led classroom observation study.

STUDY

Infer grit, motivation, and self-efficacy from chats and submissions, and feed the findings back into both the in-person program and the AI buddy.

The lab is looking for **co-owners**

FOR FUNDERS

Co-own the lab and fund what compounds: rapid 3-to-6-month studies, plus the shared data and experimentation infrastructure that makes each next study cheaper than the last.

FOR RESEARCHERS

Bring a study. TAF handles the agreements, field partners, and data pipelines. You design and run it, publish openly, and leave one reusable artifact behind.

FOR PRACTITIONERS

Bring your AI product and your questions. Get agency-centered design support and measurement you can act on, while your organization keeps full ownership of its data.

Start the conversation: **zezhen@agency.fund**

Let's build AI that
strengthens human agency,
and **prove that it does.**