


A RESEARCH AND LEARNING LAB INCUBATED AT THE AGENCY FUND

The **Agency** & AI Lab


A²L is a decentralized, virtual, and open lab with two jobs: **design** AI that strengthens human agency in the social sector, and **measure** what AI does to the agency of the people who use it.

Fifteen AI products, already on the map

Through the accelerator we run with **OpenAI** and the **Center for Global Development**, fifteen organizations built and shipped AI for the people they serve, across fourteen countries. Hover any organization to see its product and the evidence behind it.



Our worldview of agency in a flywheel



AI can **expand** agency, and it can hamper it

CAN EXPAND

THE FULL FINDING

People stay with tools that **make them feel more capable**, even when the tools make mistakes. AI can strengthen agency in five ways: instrumental, cognitive, affective, relational, and structural.

FLIP BACK

CAN HAMPER

THE FULL FINDING

A ChatGPT-style math tutor boosted practice scores by 48%, but students scored **17% below the control group** once the AI was removed. A guaranteed tutor that offered hints rather than solutions eliminated the harm, yet students using either tool remained unaware of the gap between feeling supported and actually learning.

FLIP BACK

CAN HAMPER

THE FULL FINDING

Across 1,375 participants, an accurate chatbot raised reasoning accuracy by 25 percentage points, but a faulty one drove it 15 points below unbiased performance. People followed wrong AI advice on **roughly four in five consulted trials** and became more confident even when misled; time pressure, incentives, and feedback reduced but did not eliminate the effect.

FLIP BACK

In all three studies, people looked engaged. However, engagement itself cannot tell expansion of agency from its erosion, **which is why we measure agency directly.**

Sources: [How AI Is Affecting Human Agency](#) (Jigsaw, July 2025) · [Generative AI without guardrails can harm learning](#) (PNAS 122(26), 2025) · [Thinking, Fast, Slow, and Artificial: the rise of cognitive surrender](#) (Wharton working paper, 2026).

WHAT THE ACCELERATOR TAUGHT US

Two gaps sit between these tools and lasting impact

THE DESIGN GAP

AI products are built to optimize service delivery, not to strengthen agency.

Decades of research show how psychologically wise interventions shift meaning, unlock agency, and improve development outcomes. That knowledge rarely reaches AI product design, testing, or iteration.

THE MEASUREMENT GAP

AI products are underused as windows into how agency evolves.

For AI companies, engagement is a fine proxy for revenue. In the social sector, usage alone does not produce learning, health, or income, and it can even work against them. We analyzed 395,717 conversations on one partner platform: the average lasted 3.1 turns. The signals are there, and almost nobody is measuring them.

395,717

CONVERSATIONS ANALYZED

3.1

TURNS ON AVERAGE

The Agency & AI Lab exists to close **both gaps**.

PART 2 · THE LAB

A lab **purpose-built** for the social sector.

Inspired by industry AI-for-good labs, incubated at The Agency Fund, and open for partners to co-own.

WHAT A²L IS

One lab with **two jobs**

The Agency & AI Lab (A²L) is a decentralized, virtual, and open research and learning lab. We define **agency** as a person's capability to navigate toward the future they want.

01

Design AI for human agency

Bring the science of **wise interventions** into AI products through prompting strategies, advisory knowledge bases, and evaluation frameworks, so tools support people to think, choose, and act with confidence.

FOR NGO BUILDERS AND THE RESEARCHERS WHO ADVISE THEM

02

Measure how agency evolves with AI

Track **internally experienced agency**, what people come to believe about themselves and the AI, and **externally expressed agency**, the goals people set and pursue directly, through systems, and together with others.

FROM SURVEYS, LOG DATA, AND THE CONVERSATIONS THEMSELVES

Every A²L project serves both jobs: better designs create the data that better measures need, and **better measures make better designs possible.**

Every project runs on a **tri-sector** team

TAF

The Agency Fund, the backbone

- Manages NGO relationships and executes the data, legal, and ethics agreements.
- Maintains shared data and experimentation infrastructure, so each new study starts fast and runs with minimal engineering overhead.

R

Researchers, academic and industry

- Lead time-bound field studies, typically 3 to 6 months, co-designed with TAF and the NGO.
- Leave at least one reusable artifact behind: a dataset, a measurement scale, a model, or an intervention script. Methods ship openly unless ethics or law restrict them.

NGO

NGO partners, the ground truth

- Keep projects anchored in frontline reality: recruitment, implementation, and the interpretation of findings.
- Co-own learning goals, and retain ownership of their data and their users' data.

Every study is time-boxed and leaves a reusable artifact behind, and every method ships **open by default**.

Putting the agency literature in practitioners' hands

Workshops, assistants, and toolkits that help NGO teams design for agency and measure it, built with partners and released for reuse.

Storytelling design workshop

A NotebookLM-grounded workshop that helped **Noora Health** teams design a storytelling intervention from the narrative psychology literature.

Agency Intervention Learning Assistant

A NotebookLM assistant that answers intervention design questions from TAF's curated agency research library.

Agency/AI University

A demo of AI-assisted, self-paced courses that turn TAF playbooks into bite-size modules for NGO teams.

Intervention development playbook

A step-by-step playbook for building agency-based interventions, now in progress.

Behavioral-science skills for AI assistants

Claude Skills co-developed with **Irrational Labs** that bring behavioral science into everyday AI-assisted work.

Cultural defaults for GenAI

Early work on making AI's default interaction styles fit local models of agency, building on **Hazel Markus'** cultural psychology.

■ All artifacts are open and reusable across organizations.

Four questions set the research agenda

1 Does agency-centric AI outperform directive AI?

When an assistant is explicitly designed to build goal-setting and navigation, do users gain more agency than with a purely informational tool?

2 Which conversational strategies strengthen agency?

Which micro-interventions and theory-based techniques, embedded in AI conversations, move agency the most?

3 How should humans and AI share the work?

How do hybrid programs keep AI synergizing with human support, rather than replacing it, across contexts and populations?

4 Can agency be measured from conversations?

Can it be inferred reliably from human-AI transcripts, and do its linguistic markers differ across cultures, languages, and domains?

NOW RUNNING · NOORA HEALTH

Noora Health's programs reach **2 million users in 11 languages**, with about 10,000 incoming messages a day. Since late 2025, an AI co-pilot has helped staff answer families while **humans keep control of every reply**.

We are studying whether the co-pilot changes how staff write, and how families respond, across Bangladesh, India, and Nepal.

PART 3 · MEASURING AGENCY

You cannot design for what you **cannot see.**

Agency measurement is interesting, necessary, and complex. The next four slides show the approach we teach and use.

What you need to measure also depends on the forms of agency and how agency manifests in your product.

Agency with respect to the end goals

Can users make progress toward outcomes they value, whether through, around, or beyond the tool?

Agency with respect to the AI tool

Can users understand, question, adapt, and control the tool, deciding when and how to rely on it?

Internally experienced agency

People's thoughts and feelings

ON-PLATFORM DATA • IN-APP MICRO-SURVEY

Short in-app surveys probing users' general sense of self-efficacy or perceived control over their ultimate goal over time (e.g., "How confident is the user about developing a successful business idea? How much control do they feel over reaching their ultimate goal?")

OFF-PLATFORM DATA • PHONE OR IN-PERSON SURVEY

Longer phone/in-person surveys using validated agency measures that systematically probe users' general sense of self-efficacy or perceived control over their ultimate goal over time.

FLIP BACK

ON-PLATFORM DATA • IN-APP MICRO-SURVEY

Short in-app surveys probing whether users felt able to understand, question, or adopt an AI mentor's advice.

OFF-PLATFORM DATA • VALIDATED SCALES

Longer phone/in-person surveys using validated agency measures that systematically probe users' trust in IA and perceived dependence versus empowerment when using a tool.

FLIP BACK

Externally expressed agency

People's choices and actions

ON-PLATFORM DATA • CONVERSATION LOGS

Longitudinal user signals of agency in AI conversation logs, especially those related to setting and navigating longer-term goals or demonstrating resilience despite difficulties.

OFF-PLATFORM DATA • RUBRIC-RATED PITCHES

Longer phone/in-person surveys tracking how users develop their business ideas; rubric-rated observations of ideas and pitches; or classroom observation of teamwork, ideally linked to downstream outcomes like business earnings.

FLIP BACK

ON-PLATFORM DATA • DIGITAL TRACES


Digital traces (e.g., interaction patterns or conversation logs) showing whether users understood, questioned, or adopted an AI mentor's advice.

OFF-PLATFORM DATA • OBSERVATION AND INTERVIEWS

Observation or interview data (e.g., usability studies) on tool usage for developing business ideas.

FLIP BACK

Agency measurement sits at **Level 3**



Nobody can run an RCT on every product change. **Validated agency measures at Level 3** move when the product changes, and they speak to what Level 4 cares about. That also allows teams to anchor their A/B tests on intermediate measures that are psychologically and behaviorally meaningful and relatively easy to collect.

HOW WE SUGGEST YOU DO IT • THREE STEPS THAT LOOP

Teams should specify what agency means for their specific users and contexts

1

Define agency for your local problem

What kind of agency is the tool supposed to support? **For whom? To what end?** Anchor the definition in your program's theory of change.



2

Measure with the right tools, on the right data, at the right moment


Combine **on-platform** signals, such as in-app surveys and digital traces, with **off-platform** surveys, interviews, observations, and records.



3

Test and validate your measurement tools iteratively

Test for evidence of **reliability and validity** for your context.



We put the method **in a room**

- 1** Teams of grantees, funders, and staff coded **30 synthetic AI-assistant conversations** from education, health, and agriculture.
- 2** Everyone labeled the same **binary rubric**, built on the agency flywheel: motivation, action and choice, links to real-world outcomes, and updated mental models.
- 3** **Calibrate** showed people-vs-people agreement live. Then an **LLM judge** ran the exact same rubric on the same conversations.

What the room learned

Binary questions force choices on borderline cases, and those forced choices surface the definitional debates that actually matter for measurement.

Once a rubric earns agreement between people, an AI can apply it consistently across **thousands of conversations**. That is what makes agency measurement scale.

PART 4 · THE FLAGSHIP STUDY

Now we are building **the instrument.**

A joint research program with the University of Pennsylvania, running across six NGO platforms in five countries.

An **LLM judge** for human agency

Hamsa Bastani

The Wharton School, University of Pennsylvania

Osbert Bastani

Computer and Information Science, University of Pennsylvania

Zezen (Michael) Wu

The Agency Fund

The same Penn team ran the PNAS 2025 study on AI guardrails and learning.

Digital platforms already produce conversations, decisions, and, where partners can link them, real outcomes. We design **LLM judges** that read those transcripts against a structured rubric and infer each component of agency, with scores that **admit uncertainty**.

Intentionality

Does the user form plans that serve their goals?

Forethought

Do they weigh what would happen under different courses of action?

Self-reactiveness

Do they carry plans out and adapt when new information arrives?

Self-reflectiveness

Do they look back on outcomes to improve the next decision?

| Framework grounded in Social Cognitive Theory; it extends to other definitions of agency, including the capability approach.

HOW THE MEASUREMENT WORKS

From raw transcripts to **trusted scores**

STEP 1

Transcripts

Text, voice, and video from real platform use, plus linked real-world decisions where partners have them.



STEP 2

LLM judge

Fills a structured rubric for each user and time window, guided by in-context examples tailored to the domain.



STEP 3

Calibrated scores

A learned model turns rubric answers into agency scores that carry confidence estimates.



STEP 4

Validation

Multiple experts label the same transcripts, and lightweight micro-surveys ask users directly at the most informative moments.

WHAT MAKES THIS HARD, AND HOW WE HANDLE IT

Two kinds of agency

Toward end goals, and toward the platform itself. The rubric covers both.

Agency moves over time

Scores decompose into an overall and a time-varying component.

Missing evidence

When a transcript cannot support a judgment, the score says so.

User burden

Micro-surveys run under a strict per-user budget, placed where they are worth the interruption.

THE STUDY COHORT

Six partners, three domains, five countries

HEALTH



JH Jacaranda Health · Kenya

PROMPTS, maternal-health SMS with AI triage.

3.8M+

mothers reached with 22 county governments

EDUCATION



Nova Escola · Brazil

WhatsApp AI lesson planner on the national curriculum.

150K teachers in year one



The Apprentice Project · India

TAP Buddy project-based learning for ages 9 to 17, now adding mother-tongue voice AI.



Udhyam Learning Foundation

· India

AI help bot inside a real-business entrepreneurship program.

LIVELIHOODS



Digital Green

· Kenya, Ethiopia, India

FarmerChat agronomy advisor.

1.3M farmers



Kabakoo Academies · Mali

WhatsApp AI mentor with video reflections, often in Bambara.

Every partner already runs an AI product with real users and real stakes. Together they give the study conversations in **text, voice, and video** across languages and cultures.

What each partner brings, and what we will learn together



DIGITAL GREEN
KENYA · ETHIOPIA · INDIA

PRODUCT

FarmerChat answers smallholders' agronomy questions; version 2 adds visual learning cards and profile-based personalization.

DATA

1.3 million users and hundreds of thousands of conversations, linkable to adoption and yield surveys from ongoing RCTs.

STUDY

Find agency signals among power users, track how they shift from version 1 to version 2, and test whether they predict real adoption, alongside the World Bank MAGNET agency measures Digital Green already uses.



JACARANDA HEALTH
KENYA

PRODUCT

PROMPTS supports new and expecting mothers over two-way SMS with answers, AI risk triage, and referrals into care.

DATA

Conversation logs from a service that has reached **3.8M+ mothers**, linkable to routine surveys and care-visit records.

STUDY

Read how mothers navigate care pathways, link those signals to visits, and design proactive nudges for mothers who do not yet know what to ask.



KABAKOO ACADEMIES
MALI

PRODUCT

A **WhatsApp AI mentor** guides youth upskilling; learners answer curriculum prompts with short reflection videos, often in Bambara.

DATA

About **5,000 learners, 15,000 reflections**, and 4,000 video transcripts. A waitlist RCT already showed gains of 0.14 to 0.25 SD in growth mindset.

STUDY

Mine reflections for agency signals and connect them to income and community standing years later, under strict privacy protections for at-risk youth.

What each partner brings, and what we will learn together



NOVA ESCOLA
BRAZIL

PRODUCT

A **WhatsApp lesson planner** built on 40 years of teacher content and Brazil's national curriculum, made for low-connectivity classrooms.

DATA

150,000 teachers generated half a million lesson plans in the first year, and 76,127 in February 2026 alone.

STUDY

Does co-creating plans in conversation build lasting teaching practice, or only save time? Vary copilot features and measure what teachers learn.



THE APPRENTICE PROJECT
INDIA

PRODUCT

TAP Buddy guides students ages 9 to 17 through hands-on projects in coding, art, science, and financial literacy over WhatsApp; a new voice AI reaches parents and students in Hindi, Marathi, and Punjabi.

DATA

Multimodal project submissions (text, images, video) plus the new voice transcripts from family calls.

STUDY

Measure creativity and 21st-century skills as agency signals, and learn how bot-delivered feedback changes what students attempt next.



UDHYAM LEARNING FOUNDATION
INDIA

PRODUCT

Students form teams, receive seed money, run **real businesses**, and pitch on video; an AI help bot answers program and learning questions around the clock.

DATA

Help-bot chats, idea submissions, pitch videos, and sales records, with outcome data on about **2,000 students**, plus a McGill-led classroom observation study.

STUDY

Infer grit, motivation, and self-efficacy from chats and submissions, and feed the findings back into both the in-person program and the AI buddy.

The data **belongs** to the people it describes

Data are owned by participating NGOs and the communities and individuals they serve. Research access is **de-identified, tiered, logged, and time-bound**.

Data minimization

Collect only what the approved research purpose needs.

Informed consent

Culturally appropriate consent for every participant.

Purpose limitation

No partner may repurpose data beyond the agreed scope.

Ethical review

Every project goes through ethical review and follows responsible-AI practice.

Open by default

Tools, models, and methods ship open-source, with contributors credited.

Honest narratives

No extractive, misleading, or decontextualized storytelling about the people behind the data.

AN OPEN INVITATION

The lab is looking for **co-owners**

FOR FUNDERS

Co-own the lab and fund what compounds: rapid 3-to-6-month studies, plus the shared data and experimentation infrastructure that makes each next study cheaper than the last.

FOR RESEARCHERS

Bring a study. TAF handles the agreements, field partners, and data pipelines. You design and run it, publish openly, and leave one reusable artifact behind.

FOR PRACTITIONERS

Bring your AI product and your questions. Get agency-centered design support and measurement you can act on, while your organization keeps full ownership of its data.

Start the conversation: **zezhen@agency.fund**

Let's build AI that
strengthens human agency,
and **prove that it does.**